

[Open in app](#)[Sign up](#)[Sign in](#)**Medium**

Search

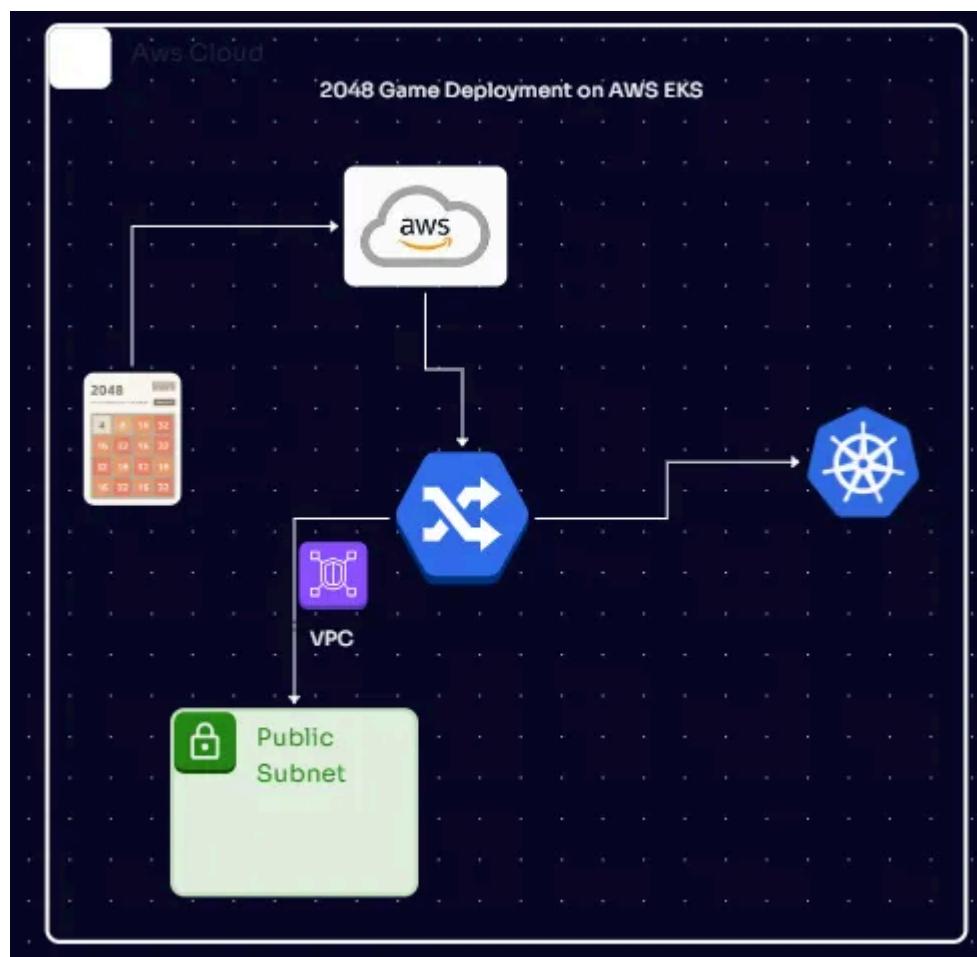


Deploying a 2048 Game on AWS EKS Fargate with Load Balancing

5 min read · Aug 25, 2024



Samsor Rahman

[Follow](#)[Listen](#)[Share](#)

Kubernetes (K8s) and Amazon Elastic Kubernetes Service (EKS) empower developers to build scalable, reliable, and containerized applications. In this guide, we'll demonstrate how to deploy the classic 2048 game on EKS, using AWS Fargate for a serverless experience and the AWS Load Balancer Controller for seamless traffic management.

Why AWS EKS Fargate & Load Balancing?

- Scalability: EKS effortlessly scales your application up or down based on demand, ensuring optimal performance during peak traffic periods.
- High Availability: AWS Load Balancer Controller automatically distributes incoming traffic across multiple pods of your application, preventing downtime even when individual components fail.
- Serverless Simplicity: Fargate eliminates the need to manage worker nodes, allowing you to focus on your application development. You pay only for the resources your application consumes, making it cost-effective.

Understanding the Components

- EKS Cluster: The foundation for our application. This is our Kubernetes environment managed by AWS.
- Fargate Profile: This enables EKS to run pods directly on Fargate, AWS's serverless compute engine.
- Namespace: A virtual cluster within Kubernetes used to organize resources. Our game will live in the game-2048 namespace.
- Deployment: Defines how the 2048 game application will be deployed (how many replicas/instances, which container image to use, etc.).
- Service: Exposes the 2048 game to other components within the Kubernetes cluster.
- Ingress: Acts as the entry point to our application from the outside world. The AWS Load Balancer Controller will create a load balancer based on this Ingress configuration.
- AWS Load Balancer Controller: This is the Kubernetes controller that integrates with AWS Application Load Balancers (ALBs). It watches for Ingress resources and provisions ALBs accordingly.
- IAM OIDC Provider: This allows EKS to authenticate with AWS IAM, enabling us to assign fine-grained permissions to our Kubernetes service accounts.

Steps for deployment

1. Create EKS Cluster

```
eksctl create cluster --name demo-cluster-1 --region us-east-1 --fargate
```

```
~/Desktop/terraform/eks
eksctl create cluster --name demo-cluster-1 --region us-east-1 --fargate

2024-06-22 16:54:25 [i] eksctl version 0.183.0-dev+a8f8fdf2e.2024-06-11T12:40:10Z
2024-06-22 16:54:25 [i] using region us-east-1
2024-06-22 16:54:27 [i] setting availability zones to [us-east-1b us-east-1f]
2024-06-22 16:54:27 [i] subnets for us-east-1b - public:192.168.0.0/19 private:192.168.64.0/19
2024-06-22 16:54:27 [i] subnets for us-east-1f - public:192.168.32.0/19 private:192.168.96.0/19
2024-06-22 16:54:27 [i] using Kubernetes version 1.30
2024-06-22 16:54:27 [i] creating EKS cluster "demo-cluster-1" in "us-east-1" region with Fargate provider
2024-06-22 16:54:27 [i] if you encounter any issues, check CloudFormation console or try 'aws cloudformation describe-stack-events --stack-name=eksctl-demo-cluster-1 --cluster=demo-cluster-1'
2024-06-22 16:54:27 [i] Kubernetes API endpoint access will use default of {publicAccess=true} for cluster "demo-cluster-1" in "us-east-1"
2024-06-22 16:54:27 [i] CloudWatch logging will not be enabled for cluster "demo-cluster-1"
2024-06-22 16:54:27 [i] you can enable it with 'eksctl utils update-cluster-logging --enable-cloudwatch-logs --region=us-east-1 --cluster=demo-cluster-1'
2024-06-22 16:54:27 [i]
2 sequential tasks: { create cluster control plane "demo-cluster-1",
  2 sequential sub-tasks: {
    wait for control plane to become ready,
    create fargate profiles,
  }
}
2024-06-22 16:54:27 [i] building cluster stack "eksctl-demo-cluster-1-cluster"
2024-06-22 16:54:30 [i] deploying stack "eksctl-demo-cluster-1-cluster"
2024-06-22 16:55:00 [i] waiting for CloudFormation stack "eksctl-demo-cluster-1-cluster"
2024-06-22 16:55:32 [i] waiting for CloudFormation stack "eksctl-demo-cluster-1-cluster"
2024-06-22 16:56:33 [i] waiting for CloudFormation stack "eksctl-demo-cluster-1-cluster"
2024-06-22 16:57:35 [i] waiting for CloudFormation stack "eksctl-demo-cluster-1-cluster"
2024-06-22 16:58:36 [i] waiting for CloudFormation stack "eksctl-demo-cluster-1-cluster"
```

This command creates an EKS cluster named “demo-cluster-1” in the us-east-1 region using Fargate.

Now Update your local kubeconfig file to connect with kubectl .

```
aws eks --region us-east-1 update-kubeconfig --name demo-cluster-1
```

This will connect your cluster with kubectl command line you will avail to control your cluster using kubectl command.

2. Create Fargate Profile and Namespace:

```
eksctl create fargateprofile \
--cluster demo-cluster-1 \
```

```
--region us-east-1 \
--name alb-sample-app \
--namespace game-2048
```

This command prepares your cluster to run the application in a serverless Fargate environment and creates a dedicated namespace (game-2048) for it.

3. Deploy the 2048 Game Application:

```
kubectl apply -f https://raw.githubusercontent.com/kubernetes-sigs/aws-load-balancer-controller/v2.0.0/deploy/managed/applications/game-2048.yaml
```

In this URL are the Deployment, Service, and Ingress resources for the 2048 game application that was already created previously.

Here is the complete breakdown of the YAML file below.

```
---
apiVersion: v1
kind: Namespace
metadata:
  name: game-2048
---
apiVersion: apps/v1
kind: Deployment
metadata:
  namespace: game-2048
  name: deployment-2048
spec:
  selector:
    matchLabels:
      app.kubernetes.io/name: app-2048
  replicas: 5
  template:
    metadata:
      labels:
        app.kubernetes.io/name: app-2048
  spec:
    containers:
      - image: public.ecr.aws/l6m2t8p7/docker-2048:latest
        imagePullPolicy: Always
        name: app-2048
        ports:
```

```
- containerPort: 80
---
apiVersion: v1
kind: Service
metadata:
  namespace: game-2048
  name: service-2048
spec:
  ports:
    - port: 80
      targetPort: 80
      protocol: TCP
    type: NodePort
  selector:
    app.kubernetes.io/name: app-2048
---
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  namespace: game-2048
  name: ingress-2048
  annotations:
    alb.ingress.kubernetes.io/scheme: internet-facing
    alb.ingress.kubernetes.io/target-type: ip
spec:
  ingressClassName: alb
  rules:
    - http:
        paths:
          - path: /
            pathType: Prefix
            backend:
              service:
                name: service-2048
                port:
                  number: 80
```

4. Configure IAM OIDC Provider and Policy:

Now Associate an IAM OIDC provider with your cluster:

```
eksctl utils associate-iam-oidc-provider --cluster demo-cluster-1 --approve
```

Then Create an IAM policy for the AWS Load Balancer Controller:

```
curl -O https://raw.githubusercontent.com/kubernetes-sigs/aws-load-balancer-controller/v2.1.0/binaries/amazonlinux2/amazonlinux2-x86_64/aws-iam-create-policy
aws iam create-policy \
--policy-name AWSLoadBalancerControllerIAMPoly \
--policy-document file://iam_policy.json
```

Now iam policy is created .

```
$ aws iam create-policy \
--policy-name AWSLoadBalancerControllerIAMPoly \
--policy-document file://iam_policy.json
{
  "Policy": {
    "PolicyName": "AWSLoadBalancerControllerIAMPoly",
    "PolicyId": "ANPARHQPWKGQN53NP7BI",
    "Arn": "arn:aws:iam::084858523853:policy/AWSLoadBalancerControllerIAMPoly",
    "Path": "/",
    "DefaultVersionId": "v1",
    "AttachmentCount": 0,
    "PermissionsBoundaryUsageCount": 0,
    "IsAttachable": true,
    "CreateDate": "2024-04-25T20:02:49+00:00",
    "UpdateDate": "2024-04-25T20:02:49+00:00"
```

Now we create the role .

```
eksctl create iamserviceaccount \
--cluster=<your-cluster-name> \
--namespace=kube-system \
--name=aws-load-balancer-controller \
--role-name AmazonEKSLoadBalancerControllerRole \
--attach-policy-arn=arn:aws:iam::<your-aws-account-id>:policy/AWSLoadBalancerControllerIAMPoly \
--approve
```

Now, I'm attaching the role to the ServiceAccount of the pod. This ensures that the pod has the necessary credentials (through the ServiceAccount) to integrate with other AWS resources.

Let's proceed with creating the Application Load Balancer (ALB) Ingress controller we discussed earlier. We'll use a Helm chart for this, which will automatically create the controller pod and the ServiceAccount associated with it.

```
helm repo add eks https://aws.github.io/eks-charts
```

```
helm repo update eks
```

Now lets install the helm charts so again while installing the helm chart you have to modify couple of thisg one is vpc id

```
helm install aws-load-balancer-controller eks/aws-load-balancer-controller \
-n kube-system \
--set clusterName=<your-cluster-name> \
--set serviceAccount.create=false \
--set serviceAccount.name=aws-load-balancer-controller \
--set region=<region> \
--set vpcId=<your-vpc-id>
```

Here i am using the helm chart for installing the load balancer controller

Now final check you have to do is you have to verify this load balancer is created and there are least two replicas of it .

```
$ kubectl get deployment -n kube-system aws-load-balancer-controller
NAME                  READY   UP-TO-DATE   AVAILABLE   AGE
aws-load-balancer-controller   2/2      2           2          2m25s
```

This confirms that your ALB Ingress controller is functioning correctly. The controller deployment creates two replicas, ensuring one runs in each Availability

Zone for redundancy. It continuously monitors for Ingress resources and ensures the ALB resources are also spread across two Availability Zones for high availability.

```
$ kubectl get ingress -n game-2048
NAME      CLASS   HOSTS   ADDRESS
ingress-2048   alb     *       k8s-game2048-ingress2-7b92b43b8b-788045392.us-ea
```

Done !!!!!!!

Access the URL and play with your app

DevOps

Cloud Computing

Cloud

Kubernetes

AWS



Follow

Written by Samsor Rahman

1.1K followers · 32 following

Building infrastructure that scales :)

No responses yet



Write a response

What are your thoughts?

More from Samsor Rahman



Written by Keith Marshall

Copyright © 2009-2013, MinGW.org Project

<http://mingw.org>

This is free software; see the product documentation or source code, for copying and redistribution conditions. There is NO WARRANTY; not even an implied WARRANTY OF MERCHANTABILITY, nor of FITNESS FOR ANY PARTICULAR PURPOSE.

This tool will guide you through the first time setup of the MinGW Installation Manager software (`mingw-get`) on your computer; additionally, it will offer you the opportunity to install some other common components of the MinGW software distribution.

After first time setup has been completed, you should invoke the MinGW Installation Manager directly, (either the CLI `mingw-get.exe` variant, or its GUI counterpart, according to your preference), when you wish to add or to remove components, or to



Samsor Rahman

How to Run a Makefile in Windows

Makefile is a special file containing information related to the program's flow, targets, and libraries. It is also called a description...

Sep 22, 2023

237

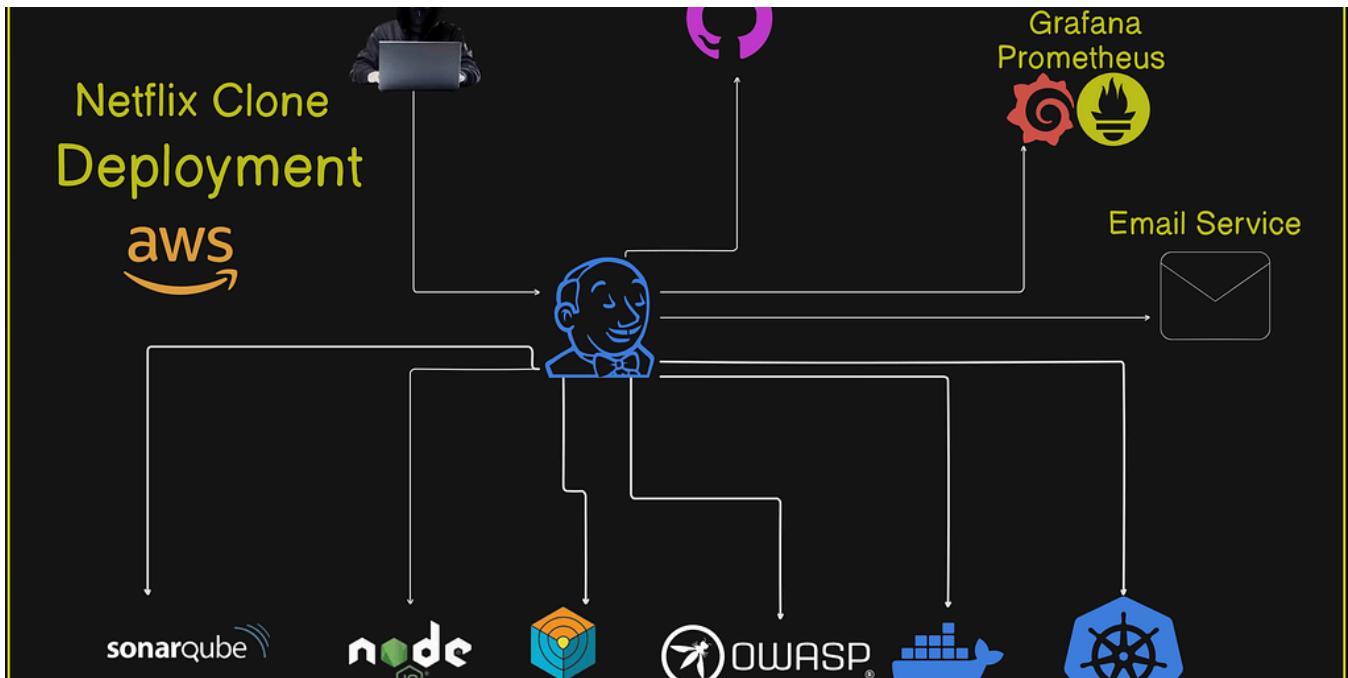
1



 Samsor Rahman

How to use Celery with Django

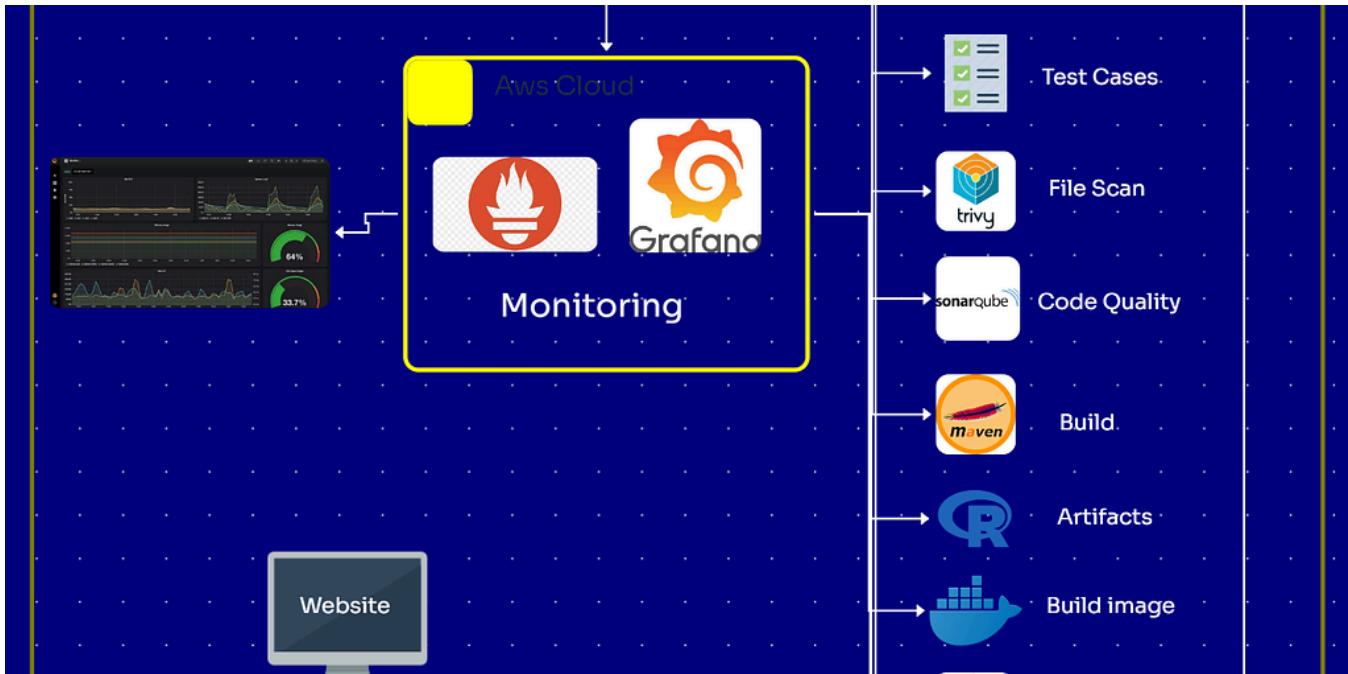
Celery is a distributed task queue system that can be used with Django to perform asynchronous tasks such as sending emails, processing...

Sep 28, 2023  458  1 In Towards AWS by Samsor Rahman

Day 89 #90 DaysOfDevOps End-to-End DevSecOps Kubernetes Project using AWS, Jenkins CICD Pipeline

Hello friends, today we will be deploying a Netflix clone! We will be using a variety of tools to automate the build, test, and deployment...

Feb 22, 2024  403  4



Samsor Rahman

Production Level CICD Project, Terraform + EKS

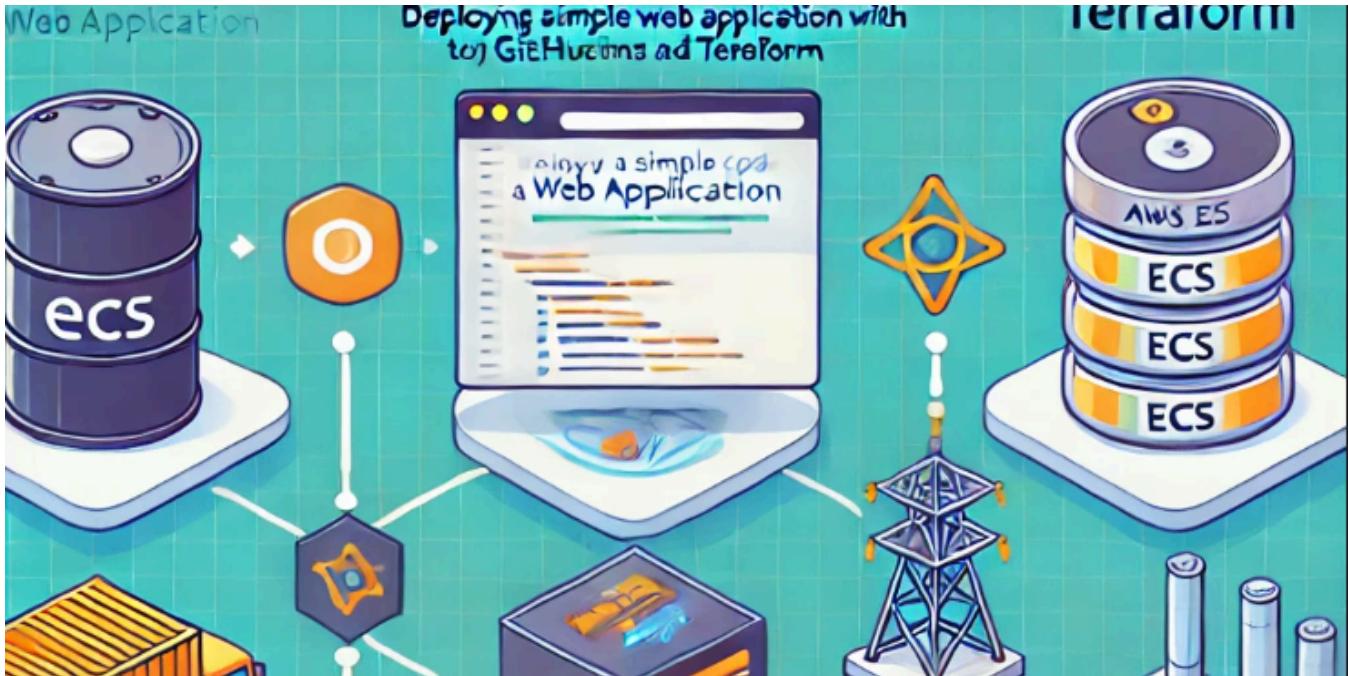
Introduction

Aug 24, 2024 144



See all from Samsor Rahman

Recommended from Medium

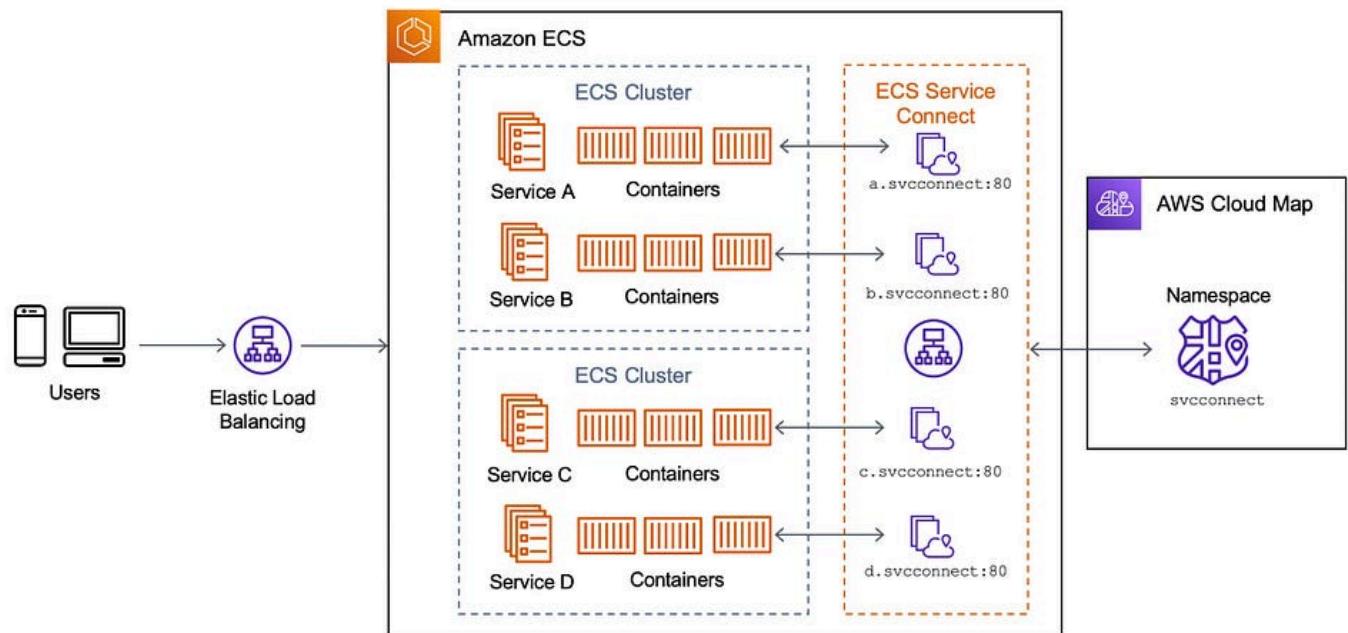


 Shrihari Haridas

Deploying a Simple Web Application to AWS ECS with GitHub Actions and Terraform

Introduction: In this blog, we will walk through the process of deploying a simple “Hello World” web application to AWS ECS (Elastic...

Feb 8 29 6

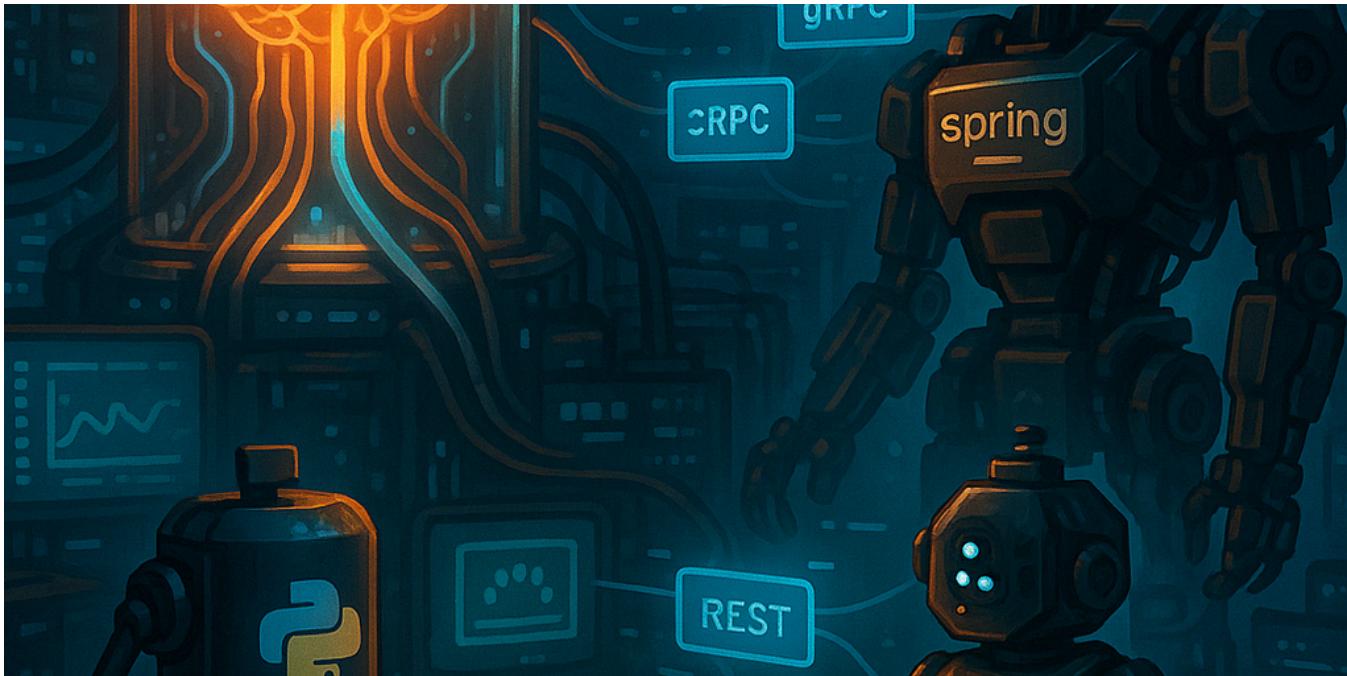


 Bhanu Reddy

ECS Best Practices— Cluster Design

1) Introduction :

⭐ Apr 2 ⚡ 10



Mr_Jeet_24

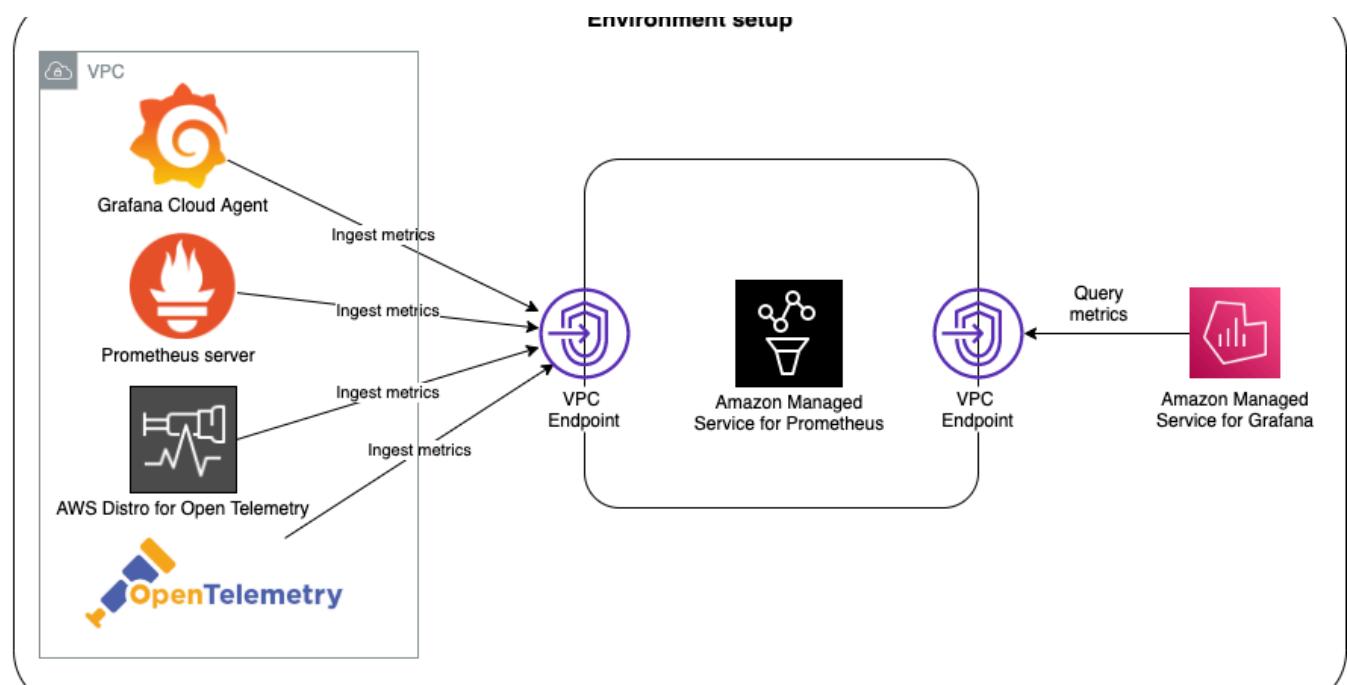
Machine Learning in Python, Deployed with Java: Architecting Hybrid AI Systems

In today's digital enterprise, it's not enough to build machine learning models—you must deploy them at scale, securely, and with high...

⭐ Jul 15 ⚡ 24



ENVIRONMENT SETUP





Nagarjun (Arjun) Nagesh

Monitoring and Logging in AWS: CloudWatch vs. Prometheus vs. ELK

Monitoring and logging are critical for ensuring application performance, security, and reliability in cloud environments. AWS offers...

★ Mar 11



Adekola Olawale

Part 16: Monitoring with Prometheus and Grafana

A Guide for DevOps Engineers

★ Apr 19

1



Secrets to Catching Production Bugs Before They Bite 🐞

How Datadog's Bits AI, Watchdog & LLM Observability Slash MTTR for Senior DevOps Engineers in 2025



In Towards AWS by Mohamed ElEmam

AI-Powered Observability: Secrets to Catching Production Bugs Before They Bite 🐞

How Datadog's Bits AI, Watchdog & LLM Observability Slash MTTR for Senior DevOps Engineers in 2025

4.5 Jul 24

16



See more recommendations